

# Final Project : Complete Project

Dhyana

2025-05-03

## Final Project

### State Information

- State Name: New Jersey
- State Abbreviation: NJ
- State FIPS Code: 34

## 1. Data Preparation

### 1.1 SAIPE Data

```
SAIPE_NJ <- read.csv("~/Downloads/SAIPE_11-04-2025.csv")
clean_data <- SAIPE_NJ |>
  select(Year, ID, Name, Poverty.Universe, Number.in.Poverty) |>
  rename(
    FIPS = ID,
    County = Name,
    Population = Poverty.Universe,
    Poverty = Number.in.Poverty
  )

clean_data <- clean_data[nchar(clean_data$FIPS) == 5, ] # Exactly 5 characters
clean_data <- clean_data[clean_data$FIPS != 34000, ] # Removing new jersey
# Total number of counties in New Jersey
clean_data |>
  distinct(FIPS, County) |>
  count()
```

```
##      n
## 1  21
```

```
# Converting to numeric
clean_data <- clean_data |>
  filter(Population != "--") |>
  mutate(
    Population = gsub(",", "", Population),
    Population = as.numeric(Population)
```

```
# Largest county
largest_county <- clean_data |>
  filter(Year == 2023) |> # 2023 is the latest year
  arrange(desc(Population)) |>
  slice(1)
largest_county |>
  select(County, Population)
```


```
##           County Population
## 1 Bergen County      948615
```

```
# 9 largest counties
top_9 <- clean_data |>
  filter(Year == 2023) |>
  arrange(desc(Population)) |>
  slice_head(n = 9)
top_9 |>
  select(County, Population)
```

```
##           County Population
## 1 Bergen County      948615
## 2 Middlesex County   841362
## 3 Essex County      828312
## 4 Hudson County     697983
## 5 Ocean County      651520
## 6 Monmouth County   636773
## 7 Union County      567554
## 8 Camden County     521373
## 9 Morris County     507188
```

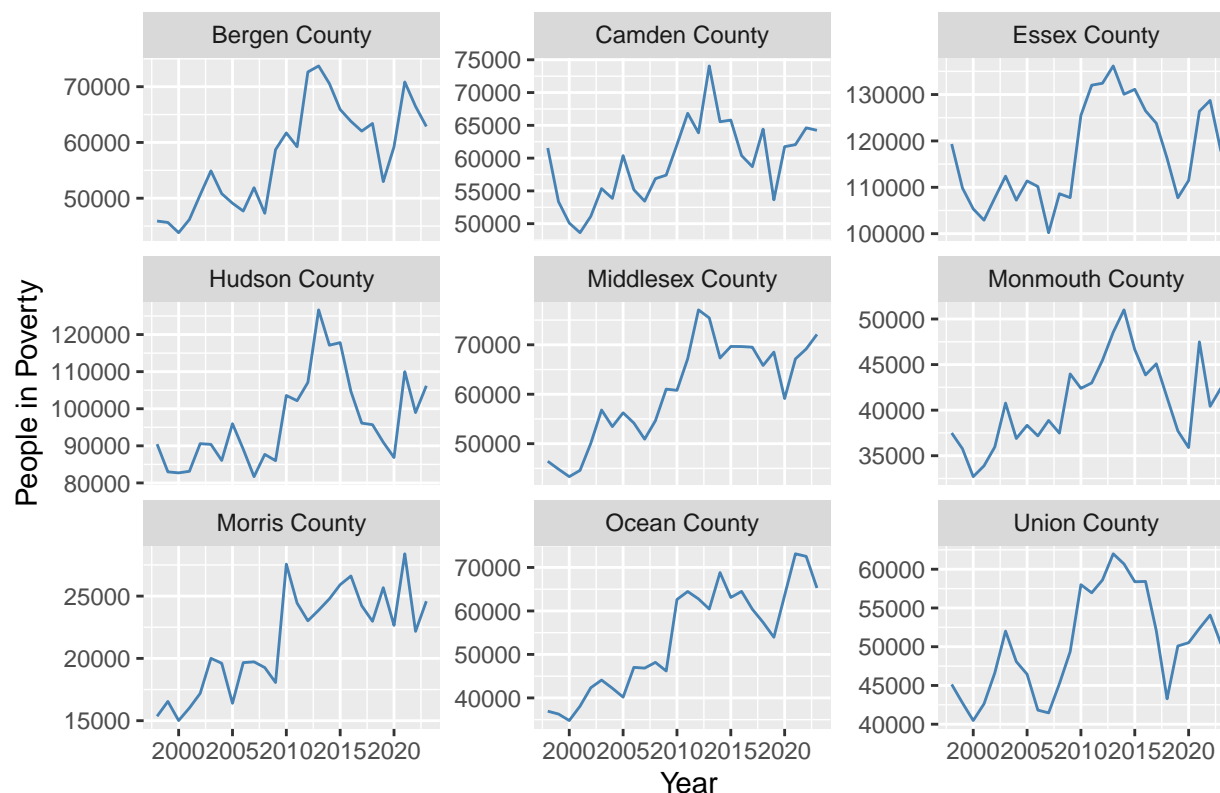
```
# Population map
data_2023 <- clean_data[clean_data$Year == 2023, ]
data_2023 <- data_2023 |>
  rename(fips = FIPS)
plot_usmap(
  data = data_2023,
  regions = "counties",
  include = "NJ",
  values = "Population"
) +
  labs(
    title = "New Jersey County Population"
  ) +
  theme(
    legend.position = "bottom",
    legend.key.width = unit(2, "cm"),
    legend.margin = margin(t = 5, unit = "pt")
  )
```

## New Jersey County Population



```
# Time Plot
clean_data$Poverty <- as.numeric(gsub(",", "", clean_data$Poverty))
top9_name <- top_9$County
top9_poverty <- clean_data[clean_data$County %in% top9_name, ]
ggplot(top9_poverty, aes(x = Year, y = Poverty)) +
  geom_line(color = "steelblue") +
  facet_wrap(~ County, scales = "free_y") +
  labs(
    title = "Poverty in top 9 counties",
    x = "Year",
    y = "People in Poverty"
  )
)
```

## Poverty in top 9 counties



## 1.2 County SNAP Benefits

```
snap_data <- read_csv("~/Creative Cloud Files/cleaned_cntysnap.csv")
```

```
## Warning: One or more parsing issues, call 'problems()' on your data frame for details,
## e.g.:
##   dat <- vroom(...)
##   problems(dat)
```

```
## Rows: 3214 Columns: 32
## -- Column specification -----
## Delimiter: ","
## chr  (6): Name, Jul-2006, Jul-2002, Jul-2001, Jul-1998, Jul-1997
## dbl  (2): State FIPS code, County FIPS code
## num (24): Jul-2022, Jul-2021, Jul-2020, Jul-2019, Jul-2018, Jul-2017, Jul-20...
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

```
snap_data <- snap_data |>
  rename(
    state_fips = `State FIPS code`,
    county_fips = `County FIPS code`,
```

```

    county_name = Name
  )
SNAP_NJ <- snap_data |>
  filter(state_fips == 34, county_fips != 0) |>
  mutate(FIPS = as.character(state_fips * 1000 + county_fips))
SNAP_NJ <- SNAP_NJ |>
  mutate(across(starts_with("Jul-"), as.numeric))

```

```

## Warning: There were 5 warnings in 'mutate()'.
## The first warning was:
## i In argument: 'across(starts_with("Jul-"), as.numeric)'.
## Caused by warning:
## ! NAs introduced by coercion
## i Run 'dplyr::last_dplyr_warnings()' to see the 4 remaining warnings.

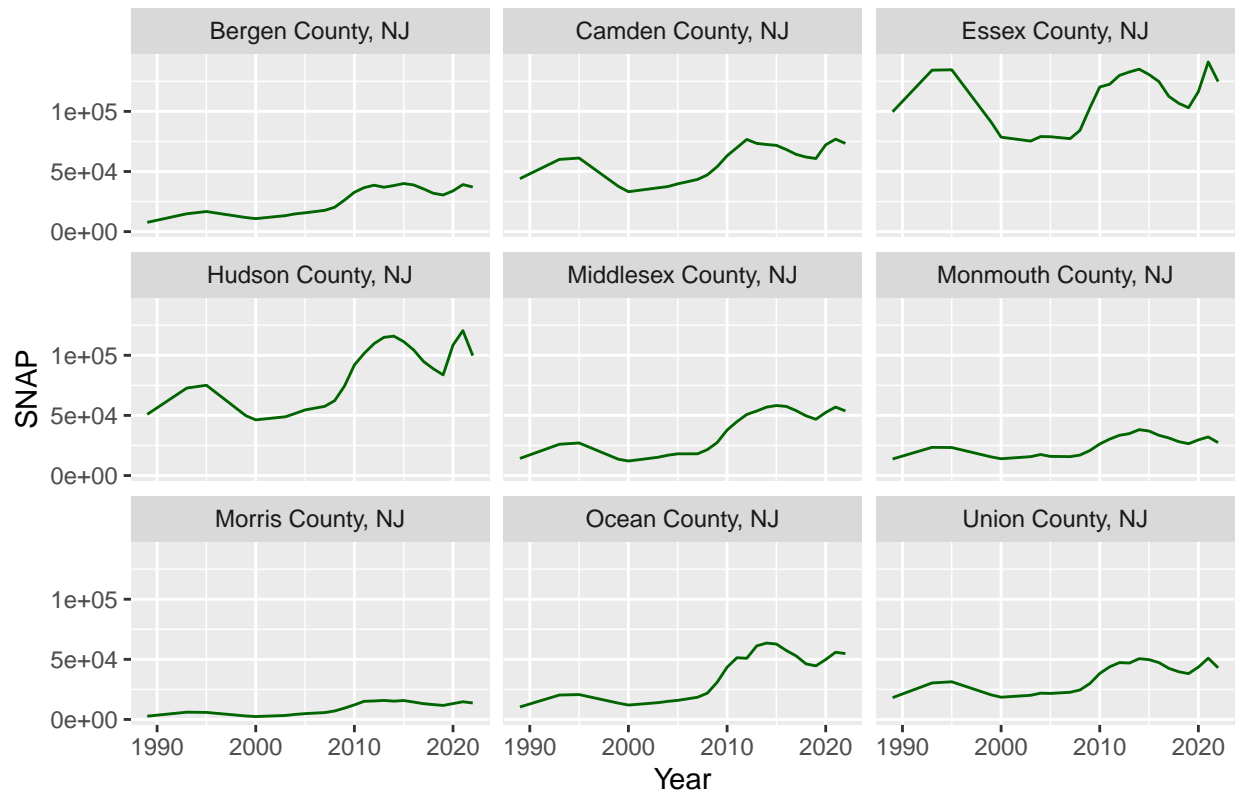
```

```

snap_long <- SNAP_NJ |>
  pivot_longer(
    cols = starts_with("Jul-"),
    names_to = "Year",
    values_to = "SNAP"
  ) |>
  mutate(Year = as.numeric(sub("Jul-", "", Year)))
top_9$FIPS <- as.character(top_9$FIPS)
snap_long$FIPS <- as.character(snap_long$FIPS)
top9_fips <- top_9$FIPS # Top 9
top9_snap <- snap_long |>
  filter(FIPS %in% top9_fips)
top9_snap <- top9_snap |>
  filter(!is.na(SNAP))
# Time PLOT
ggplot(top9_snap, aes(x = Year, y = SNAP)) +
  geom_line(color = "darkgreen") +
  facet_wrap(~ county_name) +
  labs(
    title = "SNAP Benefit Trends for Top 9 NJ Counties (by FIPS)",
    x = "Year",
    y = "SNAP"
  )

```

## SNAP Benefit Trends for Top 9 NJ Counties (by FIPS)



### 1.3 State IRS Data

```
irs_data <- read_csv("~/Creative Cloud Files/cleaned_irs.csv")
```


```
## Rows: 1734 Columns: 13
## -- Column specification -----
## Delimiter: ","
## chr  (1): Name
## dbl  (2): State FIPS code, Year
## num (10): Total exemptions, Poor exemptions, Age 65 and over exemptions, Age...
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

```
# Filtering for New Jersey
IRS_NJ <- irs_data |>
  filter(Name == "New Jersey") |>
  mutate(`Poor exemptions` = as.numeric(`Poor exemptions`)) # Converting to numeric
# Time Plot
ggplot(IRS_NJ, aes(x = Year, y = `Poor exemptions`)) +
  geom_line(color = "purple") +
  labs(
    title = "Poor Exemptions filed in New Jersey",
```

```

x = "Year",
y = "Number of Poor Exemptions"
)

```



#### 1.4 Merging the data

```

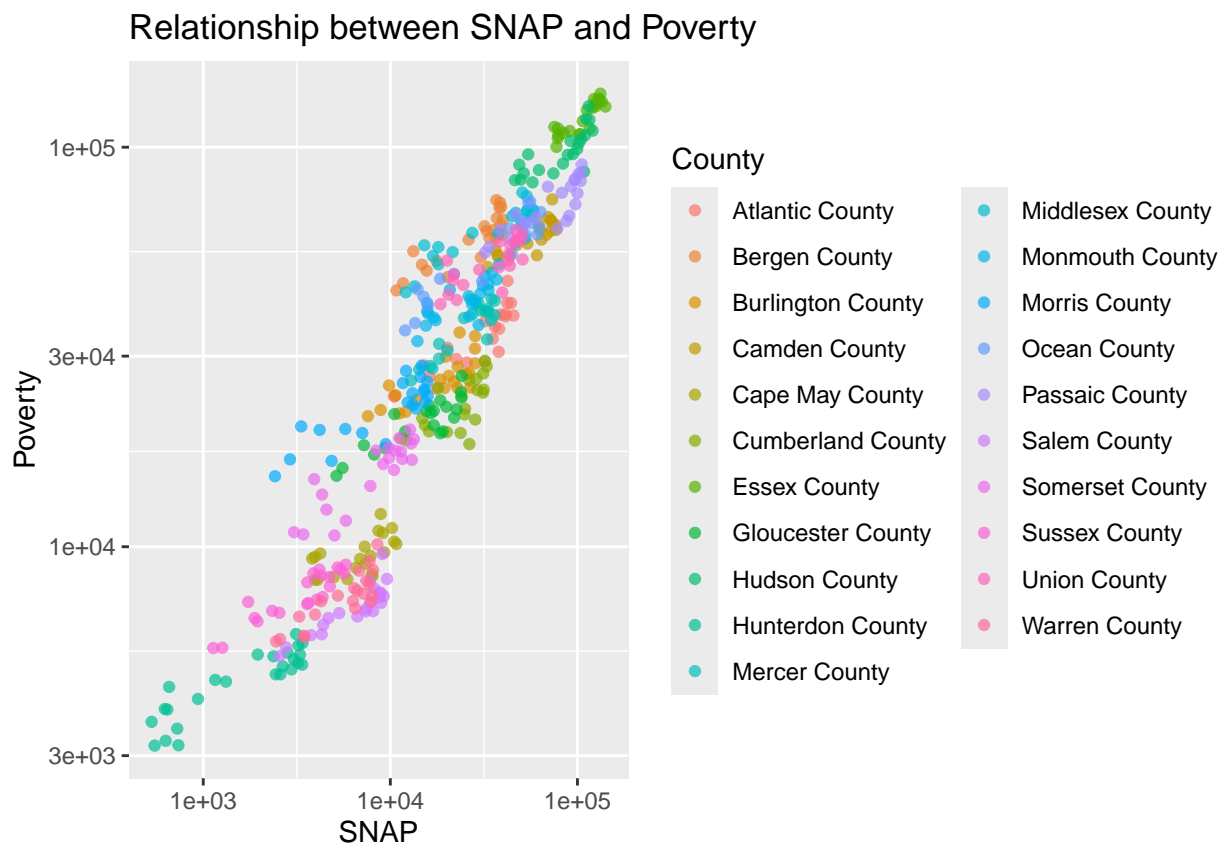
clean_data <- clean_data |>
  mutate(FIPS = as.character(FIPS))
snap_long <- snap_long |>
  mutate(FIPS = as.character(FIPS))
# Merging SAIPE and SNAP
merged_data <- clean_data |>
  filter(Year >= 1997) |>
  left_join(
    snap_long |>
      filter(Year >= 1997) |>
      select(FIPS, Year, SNAP),
    by = c("FIPS", "Year"))
# Adding IRS
merged_data <- merged_data |>
  left_join(
    IRS_NJ |>
      filter(Year >= 1997) |>

```

```

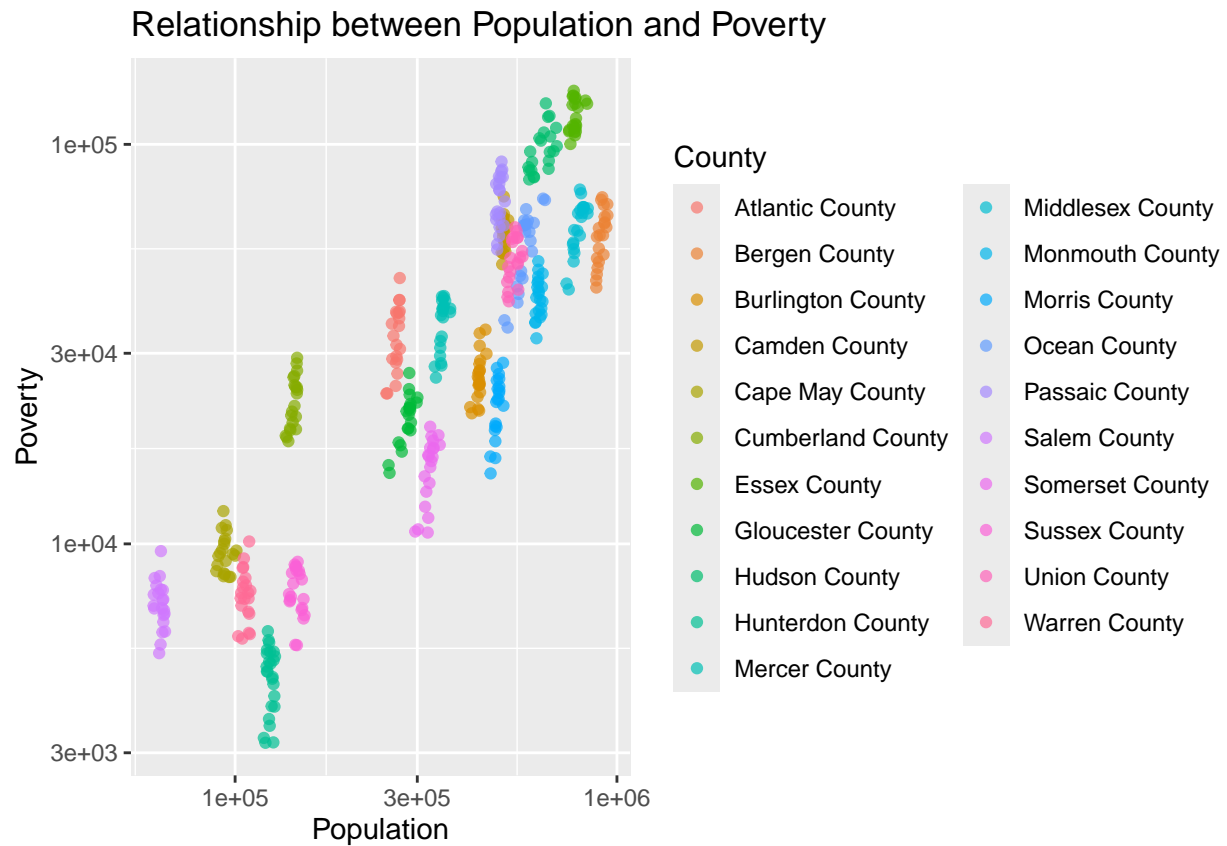
    select(Year, `Poor exemptions`),
    by = "Year"
  )
  # Removing missing values
merged_data <- merged_data |>
  filter(
    !is.na(Poverty),
    !is.na(Population),
    !is.na(SNAP),
    !is.na(`Poor exemptions`)
  )
  # Converting to a tsibble
merged_tsibble <- merged_data |>
  as_tsibble(key = c(FIPS, County), index = Year)
  # Visualization 1 - Poverty vs SNAP
ggplot(merged_tsibble, aes(x = SNAP, y = Poverty, color = County)) +
  geom_point(alpha = 0.7) +
  scale_x_log10() +
  scale_y_log10() +
  labs(
    title = "Relationship between SNAP and Poverty",
    x = "SNAP",
    y = "Poverty"
  )

```






```
# Visualization 2 - Poverty vs Population
ggplot(merged_tsibble, aes(x = Population, y = Poverty, color = County)) +
  geom_point(alpha = 0.7) +
  scale_x_log10() +
  scale_y_log10() +
  labs(
    title = "Relationship between Population and Poverty",
    x = "Population",
    y = "Poverty"
  )
)
```



```
# Visualization 3 - Poverty vs Poor Exemptions
ggplot(merged_tsibble, aes(x = `Poor exemptions`, y = Poverty, color = County)) +
  geom_point(alpha = 0.7) +
  scale_x_log10() +
  scale_y_log10() +
  labs(
    title = "Relationship between Poor Exemptions and Poverty",
    x = "Poor Exemptions",
    y = "Poverty"
  )
)
```

## Relationship between Poor Exemptions and Poverty



## 2. Linear Models

### 2.1 Variable Selection

```
final_data <- merged_tsibble
# Log-transformed variables
final_data <- final_data |>
  mutate(
    log_poverty = log(Poverty),
    log_population = log(Population),
    log_snap = log(SNAP),
    log_poor_exemptions = log(`Poor exemptions`)
  )
# 7 Linear Models
model_1 <- lm(log_poverty ~ log_population, data = final_data)
model_2 <- lm(log_poverty ~ log_snap, data = final_data)
model_3 <- lm(log_poverty ~ log_poor_exemptions, data = final_data)
model_4 <- lm(log_poverty ~ log_population + log_snap, data = final_data)
model_5 <- lm(log_poverty ~ log_population + log_poor_exemptions, data = final_data)
model_6 <- lm(log_poverty ~ log_snap + log_poor_exemptions, data = final_data)
model_7 <- lm(log_poverty ~ log_population + log_snap + log_poor_exemptions, data = final_data)
models_list <- list(model_1, model_2, model_3, model_4, model_5, model_6, model_7)
# Summary of each model
summary_1 <- glance(model_1)
```

```

summary_2 <- glance(model_2)
summary_3 <- glance(model_3)
summary_4 <- glance(model_4)
summary_5 <- glance(model_5)
summary_6 <- glance(model_6)
summary_7 <- glance(model_7)
# All summaries into a table
model_summary <- bind_rows(
  summary_1,
  summary_2,
  summary_3,
  summary_4,
  summary_5,
  summary_6,
  summary_7
)
model_summary <- model_summary |>
  mutate(model_number = 1:7) |>
  select(model_number, adj.r.squared, AIC, BIC)
print(model_summary)

```

```

## # A tibble: 7 x 4
##   model_number adj.r.squared   AIC   BIC
##         <int>         <dbl> <dbl> <dbl>
## 1             1         0.774  553.  565.
## 2             2         0.872  299.  312.
## 3             3         0.0143 1208. 1220.
## 4             4         0.971  -364. -348.
## 5             5         0.784  533.  550.
## 6             6         0.906  163.  179.
## 7             7         0.981 -537. -516.

```

```

# Best model
best_model <- model_summary |>
  filter(AIC == min(AIC))
print(paste("Best model is model number:", best_model$model_number))

```

```
## [1] "Best model is model number: 7"
```

```

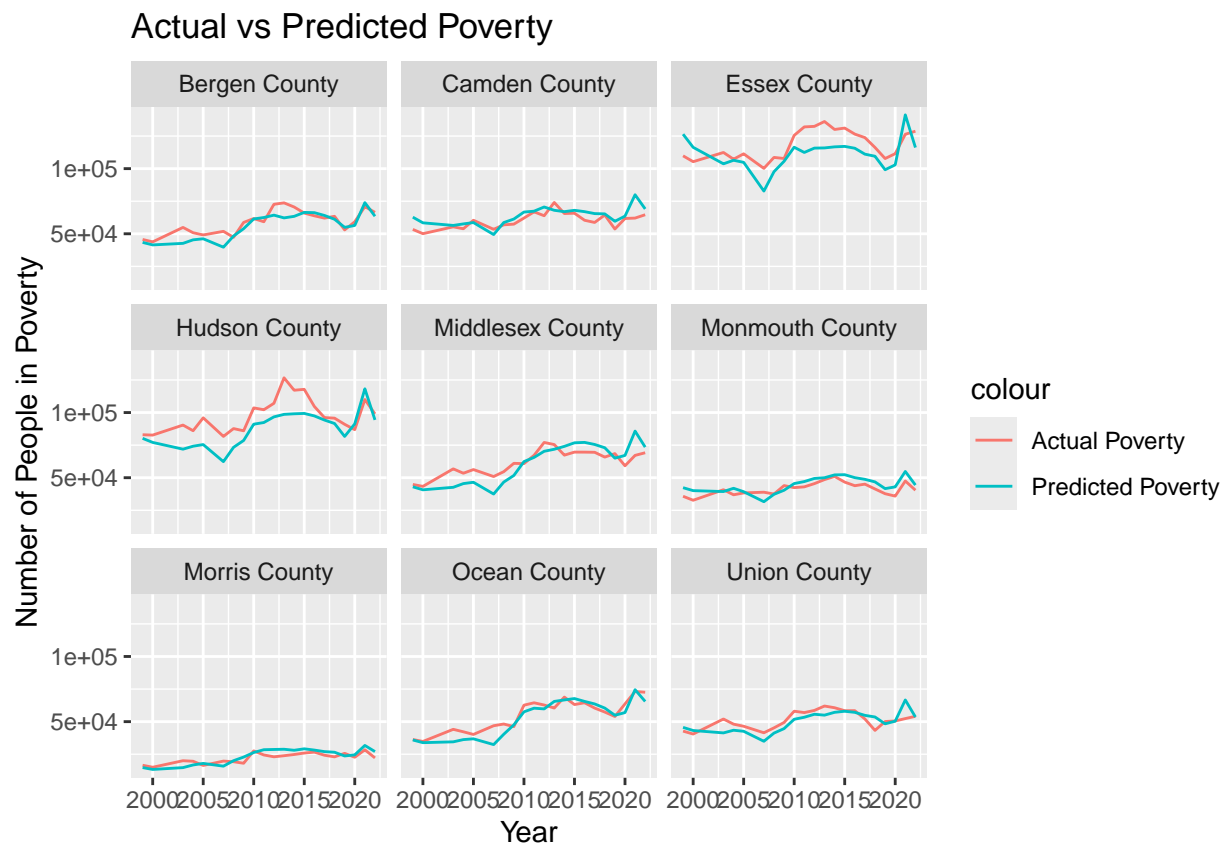
best_model <- model_7
# Predicting values
predict_values <- predict(best_model, newdata = final_data)
final_data <- final_data |>
  mutate(predict_log_poverty = predict_values)
# Converting log values back to normal
final_data <- final_data |>
  mutate(
    real_poverty = exp(log_poverty),
    predict_poverty = exp(predict_log_poverty)
  )
# Top 9 Counties
top_9_data <- final_data |>

```

```

filter(County %in% top_9$County)
# Plot
ggplot(top_9_data, aes(x = Year)) +
  geom_line(aes(y = real_poverty, color = "Actual Poverty")) +
  geom_line(aes(y = predict_poverty, color = "Predicted Poverty")) +
  facet_wrap(~County) +
  labs(
    title = "Actual vs Predicted Poverty",
    x = "Year",
    y = "Number of People in Poverty"
  )

```




## 2.2 Residual Analysis

```

# Innovation Residuals
top_9_data <- top_9_data |>
  mutate(
    innovation_residual = log_poverty - predict_log_poverty
  )
# Time Plot
ggplot(top_9_data, aes(x = Year, y = innovation_residual)) +
  geom_line(color = "purple") +
  facet_wrap(~County) +

```

```
labs(
  title = "Innovation Residuals Over Time for Top 9 Counties",
  x = "Year",
  y = "Innovation Residual"
)
```



```
# Ljung - Box Test
ljung_box <- top_9_data |>
  as_tibble() |>
  group_by(County) |>
  summarize(
    p_value = Box.test(innovation_residual, lag = 10, type = "Ljung-Box")$p.value
  )
print(ljung_box)
```

```
## # A tibble: 9 x 2
##   County      p_value
##   <chr>      <dbl>
## 1 Bergen County 0.972
## 2 Camden County 0.632
## 3 Essex County 0.412
## 4 Hudson County 0.399
## 5 Middlesex County 0.000253
## 6 Monmouth County 0.729
## 7 Morris County 0.131
```

```
## 8 Ocean County      0.0677
## 9 Union County      0.255
```

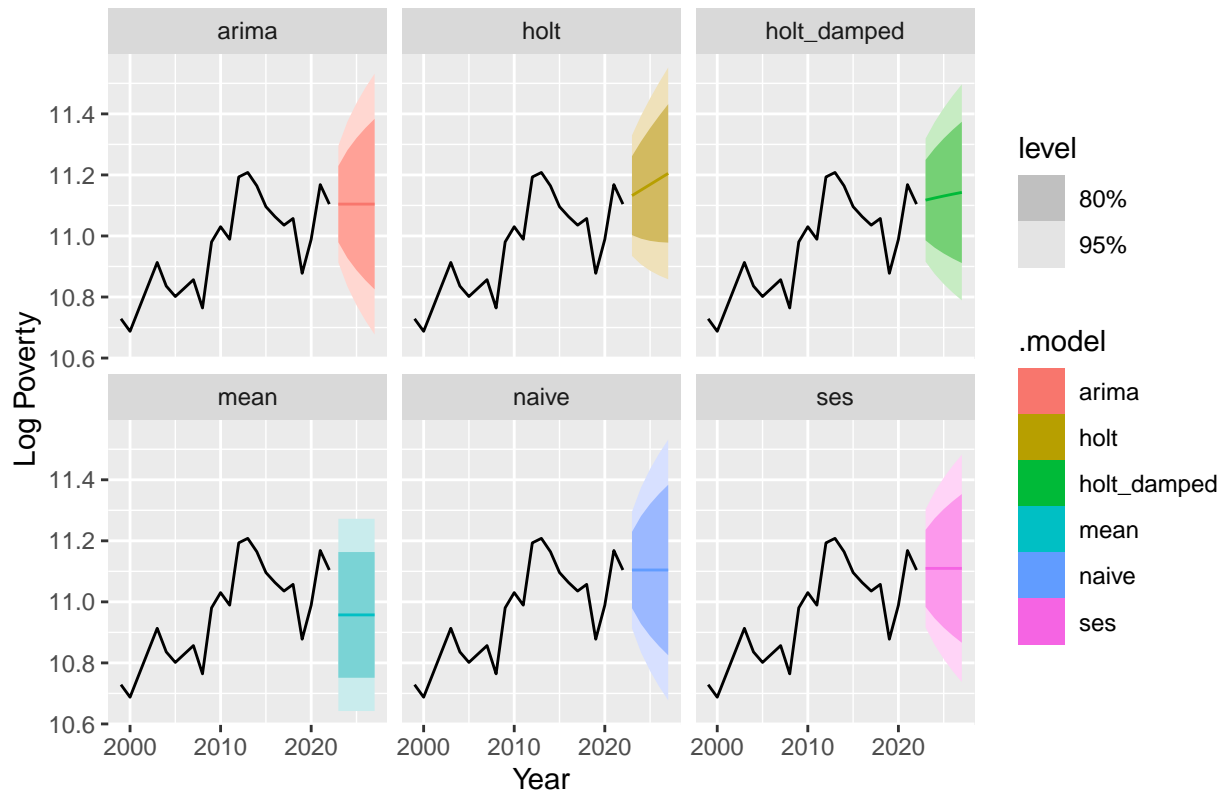
```
# Only 1 county has residuals different from white noise that is the Middlesex County
# With a high adjusted R^2 (0.981) and a strong relationship between actual and predicted values,
# the linear model performs well in predicting poverty
# According to the Ljung-Box test, innovation residuals seemed random for 8 of the 9 counties,
# indicating that the model fits the data well generally
```

### 3. Stochastic Models

#### 3.1 Single County Forecasts

```
largest_county_name <- largest_county$County[1]
county_data <- final_data |>
  filter(County == largest_county_name) |>
  select(Year, log_poverty)
largest_county_ts <- county_data |>
  as_tsibble(index = Year) |>
  tsibble::fill_gaps()
largest_county_ts <- largest_county_ts |>
  mutate(log_poverty = na.approx(log_poverty, x = Year, na.rm = FALSE))
# Models
county_models <- largest_county_ts |>
  model(
    naive = NAIVE(log_poverty),
    mean = MEAN(log_poverty),
    ses = ETS(log_poverty ~ error("A") + trend("N") + season("N")),
    holt = ETS(log_poverty ~ error("A") + trend("A") + season("N")),
    holt_damped = ETS(log_poverty ~ error("A") + trend("Ad") + season("N")),
    arima = ARIMA(log_poverty)
  )
forecasts <- county_models |>
  forecast(h = 5)
autoplot(forecasts, largest_county_ts) +
  labs(
    title = "5-Year Forecasts for Log Poverty in Largest NJ County",
    x = "Year",
    y = "Log Poverty"
  ) +
  facet_wrap(~.model)
```

## 5-Year Forecasts for Log Poverty in Largest NJ County



```
# Model Quality
model_accuracy <- accuracy(county_models)
print(model_accuracy)
```

```
## # A tibble: 6 x 10
##   .model      .type      ME    RMSE    MAE      MPE    MAPE    MASE  RMSSE    ACF1
##   <chr>      <chr>      <dbl> <dbl> <dbl>   <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 naive     Traini~  1.63e- 2 0.0975 0.0774  0.146  0.705  1    1    -0.155
## 2 mean     Traini~ -5.74e-16 0.154 0.136 -0.0198 1.24  1.76  1.58  0.744
## 3 ses      Traini~  1.95e- 2 0.0944 0.0743  0.174  0.676  0.960  0.968 -0.0222
## 4 holt     Traini~ -1.26e- 3 0.0920 0.0745 -0.0151 0.679  0.963  0.944  0.0556
## 5 holt_damped Traini~  5.18e- 5 0.0916 0.0729 -0.00399 0.663  0.941  0.939  0.0444
## 6 arima    Traini~  1.61e- 2 0.0954 0.0746  0.144  0.680  0.964  0.979 -0.154
```

```
# Among all models, the Holt's damped trend method performed best
# It has lowest RMSE and MAE
# It's forecast trend is stable, with tighter confidence intervals compared to
# ARIMA and other methods
```


## 3.2 Exponential Smoothing Models

```
final_data <- final_data |>
  mutate(log_poverty = log(Poverty))
```

```

exp_models <- final_data |>
  as_tsibble(key = c(FIPS, County), index = Year) |>
  tsibble::fill_gaps() |>
  group_by(County) |>
  mutate(log_poverty = zoo::na.approx(log_poverty, x = Year, na.rm = FALSE)) |>
  ungroup()
exp_models_fitted <- exp_models |>
  model(
    SES = ETS(log_poverty ~ error("A") + trend("N") + season("N")),
    Holt = ETS(log_poverty ~ error("A") + trend("A") + season("N")),
    Holt_Damped = ETS(log_poverty ~ error("A") + trend("Ad") + season("N"))
  )
model_accuracy <- accuracy(exp_models_fitted)
ggplot(model_accuracy, aes(x = .model, y = RMSE, fill.model)) +
  geom_col() +
  facet_wrap(~ County) +
  labs(
    title = "RMSE Comparison of Exponential Smoothing Models Across NJ Counties",
    x = "County",
    y = "RMSE",
    color = "Model"
  )

```



# I selected the Holt's damped trend model as it has the lowest RMSE in most  
 # counties, indicating the best overall forecast accuracy for poverty trends in



```
# New Jersey
```

### 3.3 ARIMA Models

```
arima_models <- final_data |>
  as_tsibble(key = c(FIPS, County), index = Year) |>
  tsibble::fill_gaps() |>
  group_by(County) |>
  mutate(log_poverty = zoo::na.approx(log_poverty, x = Year, na.rm = FALSE)) |>
  ungroup() |>
  model(auto_arima = ARIMA(log_poverty))
# ARIMA structure for each county
for (i in 1:nrow(arima_models)) {
  cat("County:", arima_models$County[i], "\n")
  print(report(arima_models$auto_arima[[i]]))
  cat("\n")
}
```

```
## County: Atlantic County
## Series: log_poverty
## Model: ARIMA(0,1,0)
##
## sigma^2 estimated as 0.01531: log likelihood=15.43
## AIC=-28.86 AICc=-28.67 BIC=-27.73
## NULL
##
## County: Bergen County
## Series: log_poverty
## Model: ARIMA(0,1,0)
##
## sigma^2 estimated as 0.009504: log likelihood=20.91
## AIC=-39.83 AICc=-39.64 BIC=-38.69
## NULL
##
## County: Burlington County
## Series: log_poverty
## Model: ARIMA(0,1,0)
##
## sigma^2 estimated as 0.01096: log likelihood=19.28
## AIC=-36.56 AICc=-36.37 BIC=-35.42
## NULL
##
## County: Camden County
## Series: log_poverty
## Model: ARIMA(1,1,0)
##
## Coefficients:
##          ar1
##        -0.4810
## s.e.    0.1788
##
```

```

## sigma^2 estimated as 0.005338: log likelihood=27.93
## AIC=-51.87 AICc=-51.27 BIC=-49.6
## NULL
##
## County: Cape May County
## Series: log_poverty
## Model: ARIMA(1,0,0) w/ mean
##
## Coefficients:
##          ar1  constant
##      0.6039    3.6242
## s.e.  0.1535    0.0155
##
## sigma^2 estimated as 0.007112: log likelihood=26.11
## AIC=-46.23 AICc=-45.03 BIC=-42.7
## NULL
##
## County: Cumberland County
## Series: log_poverty
## Model: ARIMA(1,0,0) w/ mean
##
## Coefficients:
##          ar1  constant
##      0.7138    2.8601
## s.e.  0.1383    0.0186
##
## sigma^2 estimated as 0.01077: log likelihood=21.01
## AIC=-36.01 AICc=-34.81 BIC=-32.48
## NULL
##
## County: Essex County
## Series: log_poverty
## Model: ARIMA(1,0,0) w/ mean
##
## Coefficients:
##          ar1  constant
##      0.7961    2.3804
## s.e.  0.1157    0.0097
##
## sigma^2 estimated as 0.003238: log likelihood=35.28
## AIC=-64.56 AICc=-63.36 BIC=-61.02
## NULL
##
## County: Gloucester County
## Series: log_poverty
## Model: ARIMA(0,1,0)
##
## sigma^2 estimated as 0.008564: log likelihood=22.11
## AIC=-42.22 AICc=-42.03 BIC=-41.09
## NULL
##
## County: Hudson County
## Series: log_poverty
## Model: ARIMA(1,0,0) w/ mean

```

```

##
## Coefficients:
##          ar1  constant
##          0.7020    3.4148
## s.e.    0.1398    0.0160
##
## sigma^2 estimated as 0.007986:  log likelihood=24.61
## AIC=-43.22  AICc=-42.02  BIC=-39.69
## NULL
##
## County: Hunterdon County
## Series: log_poverty
## Model: ARIMA(0,1,0)
##
## sigma^2 estimated as 0.0128:  log likelihood=18.25
## AIC=-34.49  AICc=-34.31  BIC=-33.31
## NULL
##
## County: Mercer County
## Series: log_poverty
## Model: ARIMA(0,1,1)
##
## Coefficients:
##          ma1
##          -0.3123
## s.e.    0.1665
##
## sigma^2 estimated as 0.006618:  log likelihood=25.54
## AIC=-47.08  AICc=-46.48  BIC=-44.81
## NULL
##
## County: Middlesex County
## Series: log_poverty
## Model: ARIMA(0,1,0)
##
## sigma^2 estimated as 0.006126:  log likelihood=25.97
## AIC=-49.94  AICc=-49.75  BIC=-48.8
## NULL
##
## County: Monmouth County
## Series: log_poverty
## Model: ARIMA(1,0,0) w/ mean
##
## Coefficients:
##          ar1  constant
##          0.6350    3.8718
## s.e.    0.1536    0.0165
##
## sigma^2 estimated as 0.008036:  log likelihood=24.62
## AIC=-43.24  AICc=-42.04  BIC=-39.7
## NULL
##
## County: Morris County
## Series: log_poverty

```

```

## Model: ARIMA(0,1,1) w/ drift
##
## Coefficients:
##          ma1  constant
##        -0.7347    0.0197
## s.e.    0.2199    0.0074
##
## sigma^2 estimated as 0.01418:  log likelihood=16.97
## AIC=-27.94  AICc=-26.67  BIC=-24.53
## NULL
##
## County: Ocean County
## Series: log_poverty
## Model: ARIMA(0,1,0)
##
## sigma^2 estimated as 0.009453:  log likelihood=20.98
## AIC=-39.95  AICc=-39.76  BIC=-38.82
## NULL
##
## County: Passaic County
## Series: log_poverty
## Model: ARIMA(0,1,0)
##
## sigma^2 estimated as 0.009788:  log likelihood=20.58
## AIC=-39.15  AICc=-38.96  BIC=-38.02
## NULL
##
## County: Salem County
## Series: log_poverty
## Model: ARIMA(0,1,1)
##
## Coefficients:
##          ma1
##        -0.3638
## s.e.    0.1712
##
## sigma^2 estimated as 0.007441:  log likelihood=24.17
## AIC=-44.34  AICc=-43.74  BIC=-42.07
## NULL
##
## County: Somerset County
## Series: log_poverty
## Model: ARIMA(0,1,0)
##
## sigma^2 estimated as 0.01085:  log likelihood=19.39
## AIC=-36.78  AICc=-36.59  BIC=-35.64
## NULL
##
## County: Sussex County
## Series: log_poverty
## Model: ARIMA(0,1,0)
##
## sigma^2 estimated as 0.005761:  log likelihood=26.67
## AIC=-51.34  AICc=-51.15  BIC=-50.21

```

```
## NULL
##
## County: Union County
## Series: log_poverty
## Model: ARIMA(1,0,0) w/ mean
##
## Coefficients:
##          ar1  constant
##          0.7877    2.2956
## s.e.    0.1201    0.0140
##
## sigma^2 estimated as 0.006625:  log likelihood=26.71
## AIC=-47.42  AICc=-46.22  BIC=-43.88
## NULL
##
## County: Warren County
## Series: log_poverty
## Model: ARIMA(0,1,1)
##
## Coefficients:
##          ma1
##          -0.4713
## s.e.    0.1805
##
## sigma^2 estimated as 0.01182:  log likelihood=18.79
## AIC=-33.58  AICc=-32.98  BIC=-31.31
## NULL
```

```
# The most commonly selected models by auto ARIMA across New Jersey counties are
# ARIMA(0,1,0) = selected for 9 counties
# ARIMA(1,0,0) with mean = selected for 6 counties
# ARIMA(0,1,1) and variations = selected for 4 counties
# 3 ARIMA Models
```

```
common_models_data <- final_data |>
  as_tsibble(key = c(FIPS, County), index = Year) |>
  tsibble::fill_gaps() |>
  group_by(County) |>
  mutate(log_poverty = zoo::na.approx(log_poverty, x = Year, na.rm = FALSE)) |>
  ungroup()
common_models <- common_models_data |>
  model(
    ARIMA_010 = ARIMA(log_poverty ~ pdq(0,1,0)),
    ARIMA_100_mean = ARIMA(log_poverty ~ pdq(1,0,0)),
    ARIMA_011 = ARIMA(log_poverty ~ pdq(0,1,1))
  )
# Model Quality
comparison <- accuracy(common_models)
comparison_summary <- comparison |>
  group_by(.model) |>
  summarise(mean_RMSE = mean(RMSE, na.rm = TRUE)) |>
  arrange(mean_RMSE)
print(comparison_summary)
```

```
## # A tibble: 3 x 2
```

```
##   .model      mean_RMSE
##   <chr>      <dbl>
## 1 ARIMA_011    0.0904
## 2 ARIMA_100_mean 0.0908
## 3 ARIMA_010    0.0945
```

```
# ARIMA(0,1,1) is the best model because it has the lowest RMSE
```

## 3.4 Cross Validation

```
cross_validation_data <- final_data |>
  as_tsibble(key = c(FIPS, County), index = Year) |>
  tsibble::fill_gaps() |>
  group_by(County) |>
  filter(!is.na(Poverty)) |>
  mutate(log_poverty = log(Poverty)) |>
  filter(!is.infinite(log_poverty)) |>
  ungroup() |>
  stretch_tsibble(.init = 10, .step = 1)
cross_validation_model <- cross_validation_data |>
  model(
    ETS = ETS(log_poverty ~ error("A") + trend("Ad") + season("N")),
    ARIMA = ARIMA(log_poverty ~ pdq(1,0,0))
  )
```

```
## Warning: 240 errors (1 unique) encountered for ETS
```

```
## [240] .data contains implicit gaps in time. You should check your data and convert implicit gaps into explicit gaps
```

```
## Warning: 240 errors (1 unique) encountered for ARIMA
```

```
## [240] .data contains implicit gaps in time. You should check your data and convert implicit gaps into explicit gaps
```

```
# Forecast 5 years
```

```
cross_validation_forecast <- cross_validation_model |>
  forecast(h = 5)
cross_validation_accuracy_data <- final_data |>
  as_tsibble(key = c(FIPS, County), index = Year)
cross_validation_accuracy <- cross_validation_forecast |>
  accuracy(cross_validation_accuracy_data)
```

```
## Warning: The future dataset is incomplete, incomplete out-of-sample data will be treated as missing.
```

```
## 5 observations are missing between 2023 and 2027
```

```
rmse_model <- cross_validation_accuracy |>
  group_by(.model) |>
  summarise(mean_RMSE = mean(RMSE, na.rm = TRUE)) |>
  arrange(mean_RMSE)
print(rmse_model)
```

```
## # A tibble: 2 x 2
##   .model mean_RMSE
##   <chr>      <dbl>
## 1 ETS        0.144
## 2 ARIMA      0.177
```

```
# ETS is the better choice for statewide poverty forecasting as it has the
# lowest RMSE that is 0.144
```

## 4. Forecasts

```
forecast_data <- final_data |>
  as_tsibble(key = c(FIPS, County), index = Year) |>
  tsibble::fill_gaps() |>
  group_by(County) |>
  mutate(
    log_poverty = log(Poverty),
    log_poverty = na.approx(log_poverty, x = Year, na.rm = FALSE)
  ) |>
  ungroup() |>
  filter(!is.infinite(log_poverty))
ets_forecast_model <- forecast_data |>
  model(ETS = ETS(log_poverty ~ error("A") + trend("Ad") + season("N")))
ets_forecast <- ets_forecast_model |>
  forecast(h = 5)
latest_year <- max(forecast_data$Year)
forecast_2028 <- ets_forecast |>
  filter(Year == latest_year + 5) |>
  as_tibble() |>
  mutate(predicted_poverty = exp(.mean)) |>
  select(FIPS, County, predicted_poverty)
current_data <- final_data |>
  filter(Year == latest_year) |>
  select(FIPS, County, Poverty, Population)
poverty_change <- forecast_2028 |>
  left_join(current_data, by = c("FIPS", "County")) |>
  mutate(
    increase = predicted_poverty - Poverty,
    percent_increase = 100 * (increase / Population)
  )
# Top 5 Counties
top_5_counties <- poverty_change |>
  arrange(desc(percent_increase)) |>
  slice_head(n = 5)
print(top_5_counties |> select(County, percent_increase))
```

```
## # A tibble: 5 x 2
##   County      percent_increase
##   <chr>      <dbl>
## 1 Atlantic County      1.49
## 2 Ocean County        1.38
```

```
## 3 Passaic County      1.21
## 4 Warren County      0.868
## 5 Morris County      0.859
```

```
# Map
map_data <- poverty_change |>
  mutate(fips = FIPS) |>
  select(fips, percent_increase)
# Plot
plot_usmap(
  data = map_data,
  regions = "counties",
  include = "NJ",
  values = "percent_increase"
) +
  labs(
    title = "Forecasted 5-Year % Increase in Poverty by NJ County",
    fill = "% Increase"
  )
```

Forecasted 5-Year % Increase in Poverty by NJ County

